

1. Case Study: Efficient File Synchronization in a Multi-Process Linux System

Problem Statement:

A company maintains a large number of log files that are **continuously updated** by multiple processes. To ensure data consistency, a synchronization mechanism is required that allows multiple processes to write to these log files without data corruption while maintaining high performance.

The challenge is to design and analyze an efficient file synchronization mechanism that balances data integrity, concurrency, and performance.

Requirements:

1. **Concurrent Writes:** Multiple processes should be able to **write to the same file** without **corrupting data**.
2. **Data Integrity:** Ensuring that log entries are not interleaved or lost.
3. **Performance:** The solution should minimize I/O bottlenecks and context switching overhead.
4. **Scalability:** The system should scale efficiently as the number of processes increases.

•5) **Considerations:**

Synchronization Mechanisms:

- File locking (fcntl, flock)
- Memory-mapped files (mmap)
- POSIX message queues for logging
- Named pipes (FIFOs)
- Buffering vs. Direct Writes:** Should logs be buffered before writing or written directly to disk?
- Handling Failures:** What happens if a process crashes while holding a lock?
- Performance Analysis:** Compare different approaches using system tools (strace, perf, iotop).

1.Deliverables for Analysis:


Implementation Approach: Describe different synchronization strategies and justify the best choice.

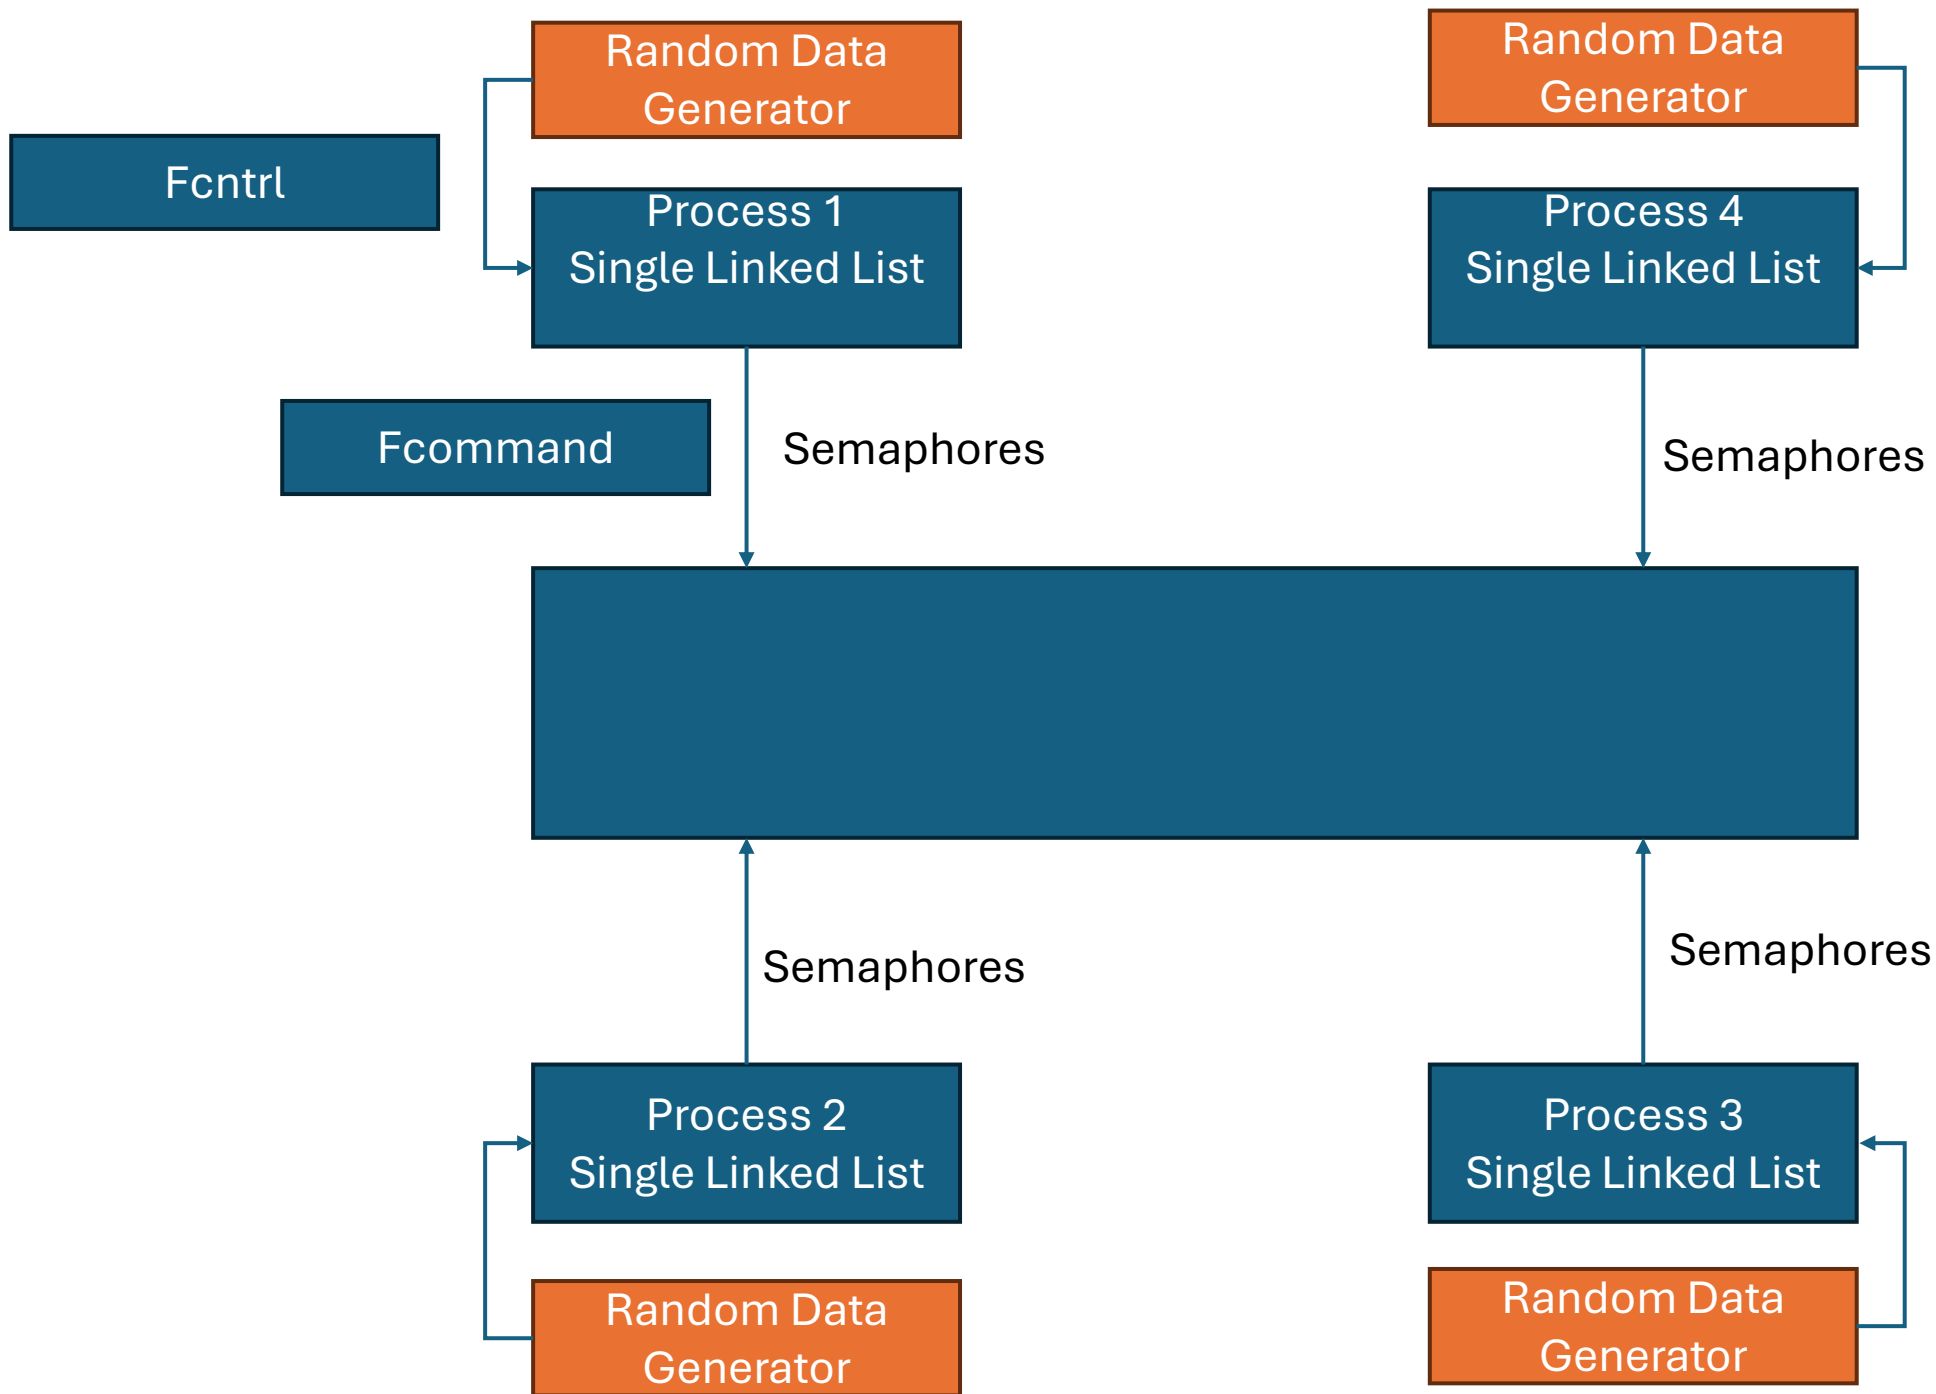
2.Performance Evaluation: Analyze the trade-offs using benchmarks.

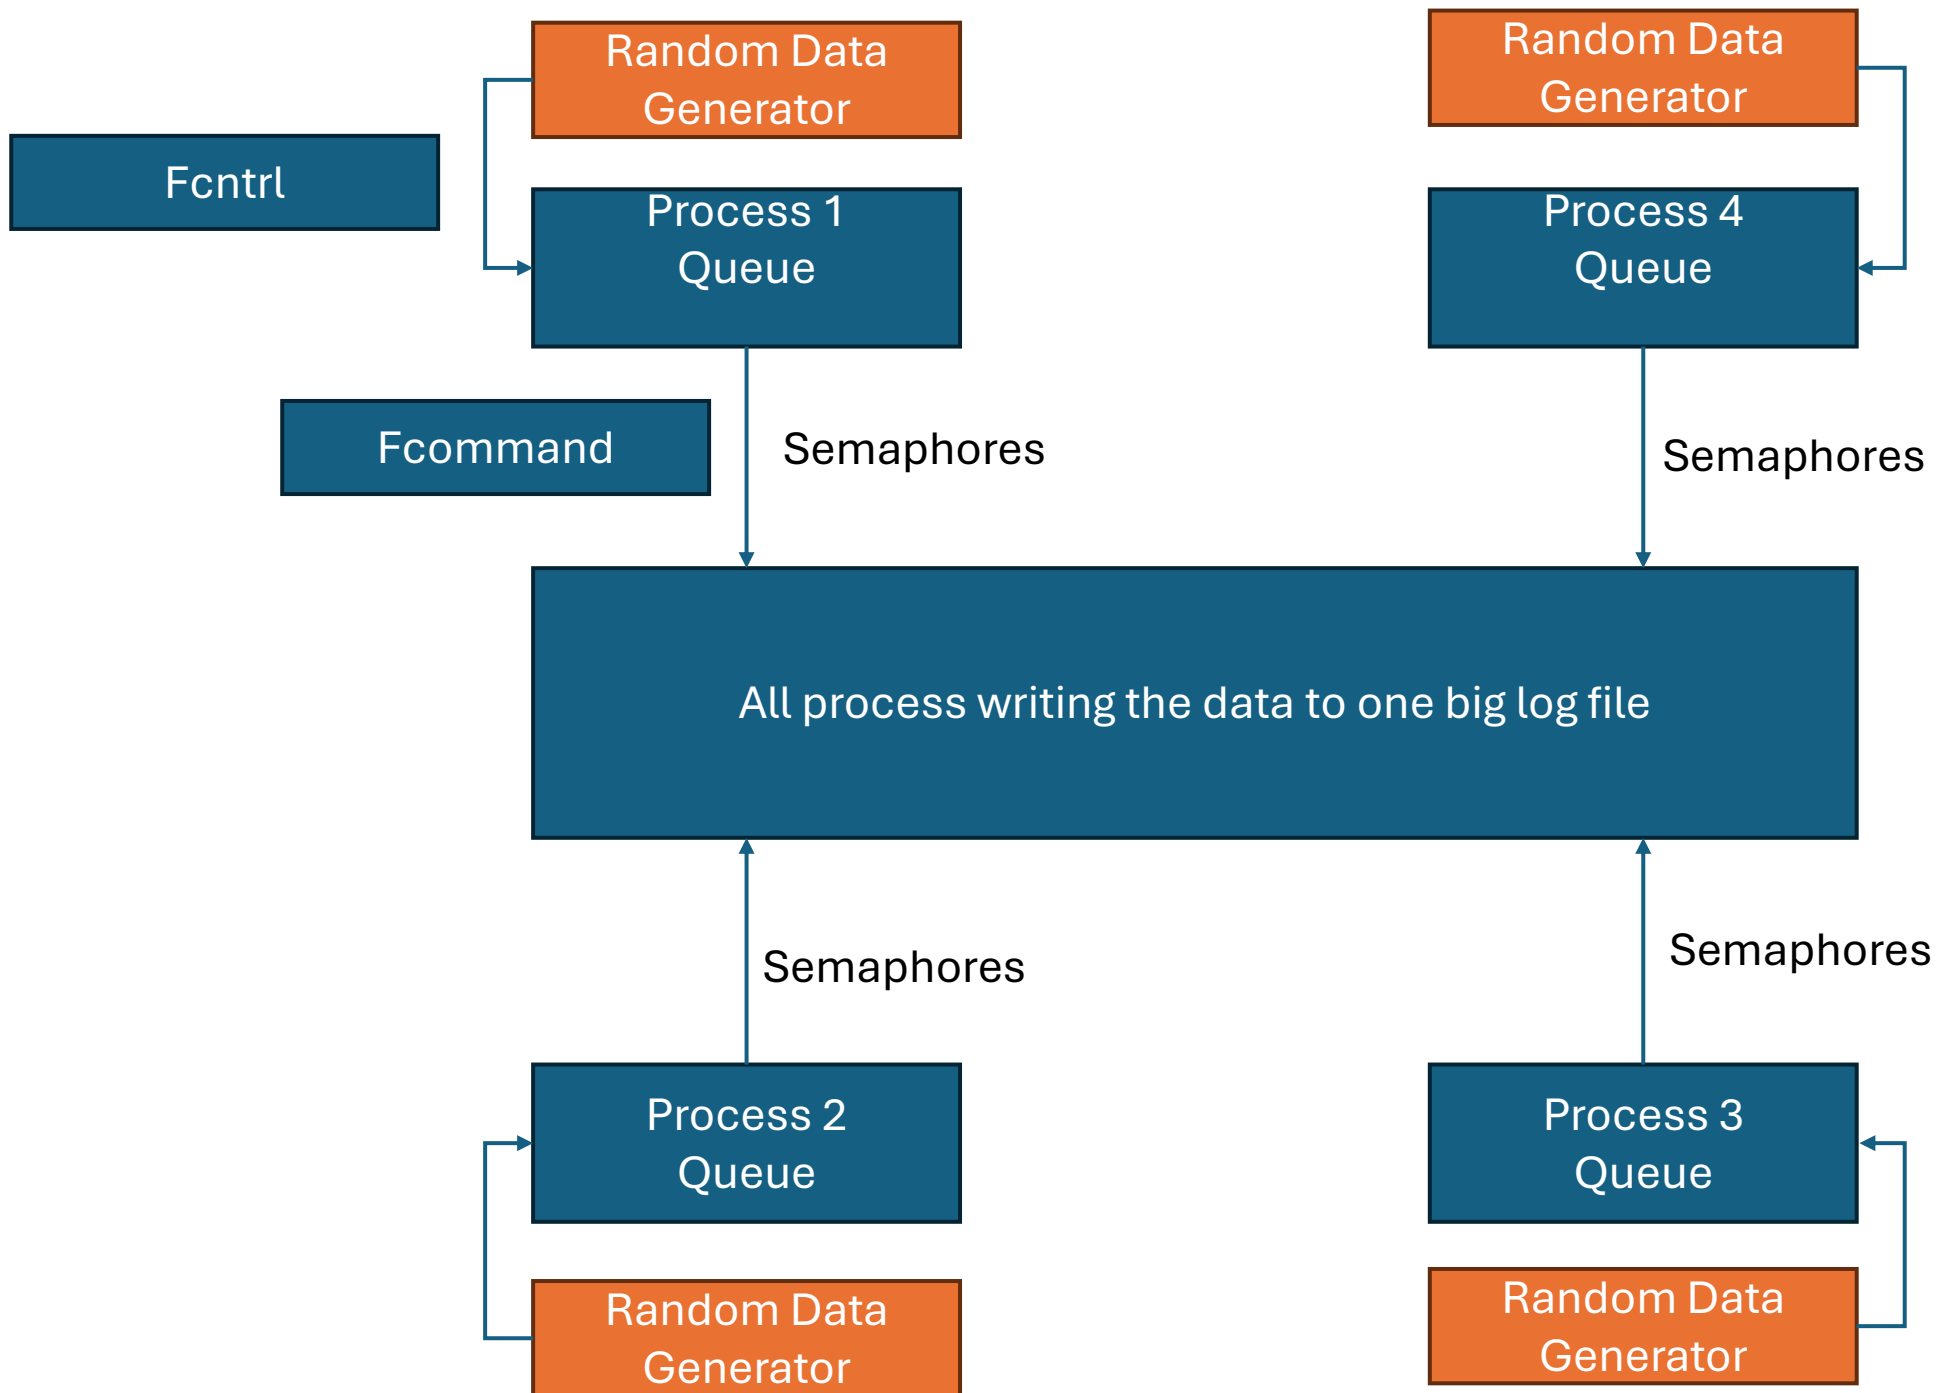
3.Failure Handling: Discuss how to ensure robustness against crashes.

4.Code Sample: Implement a prototype using C with proper error handling and logging mechanisms.

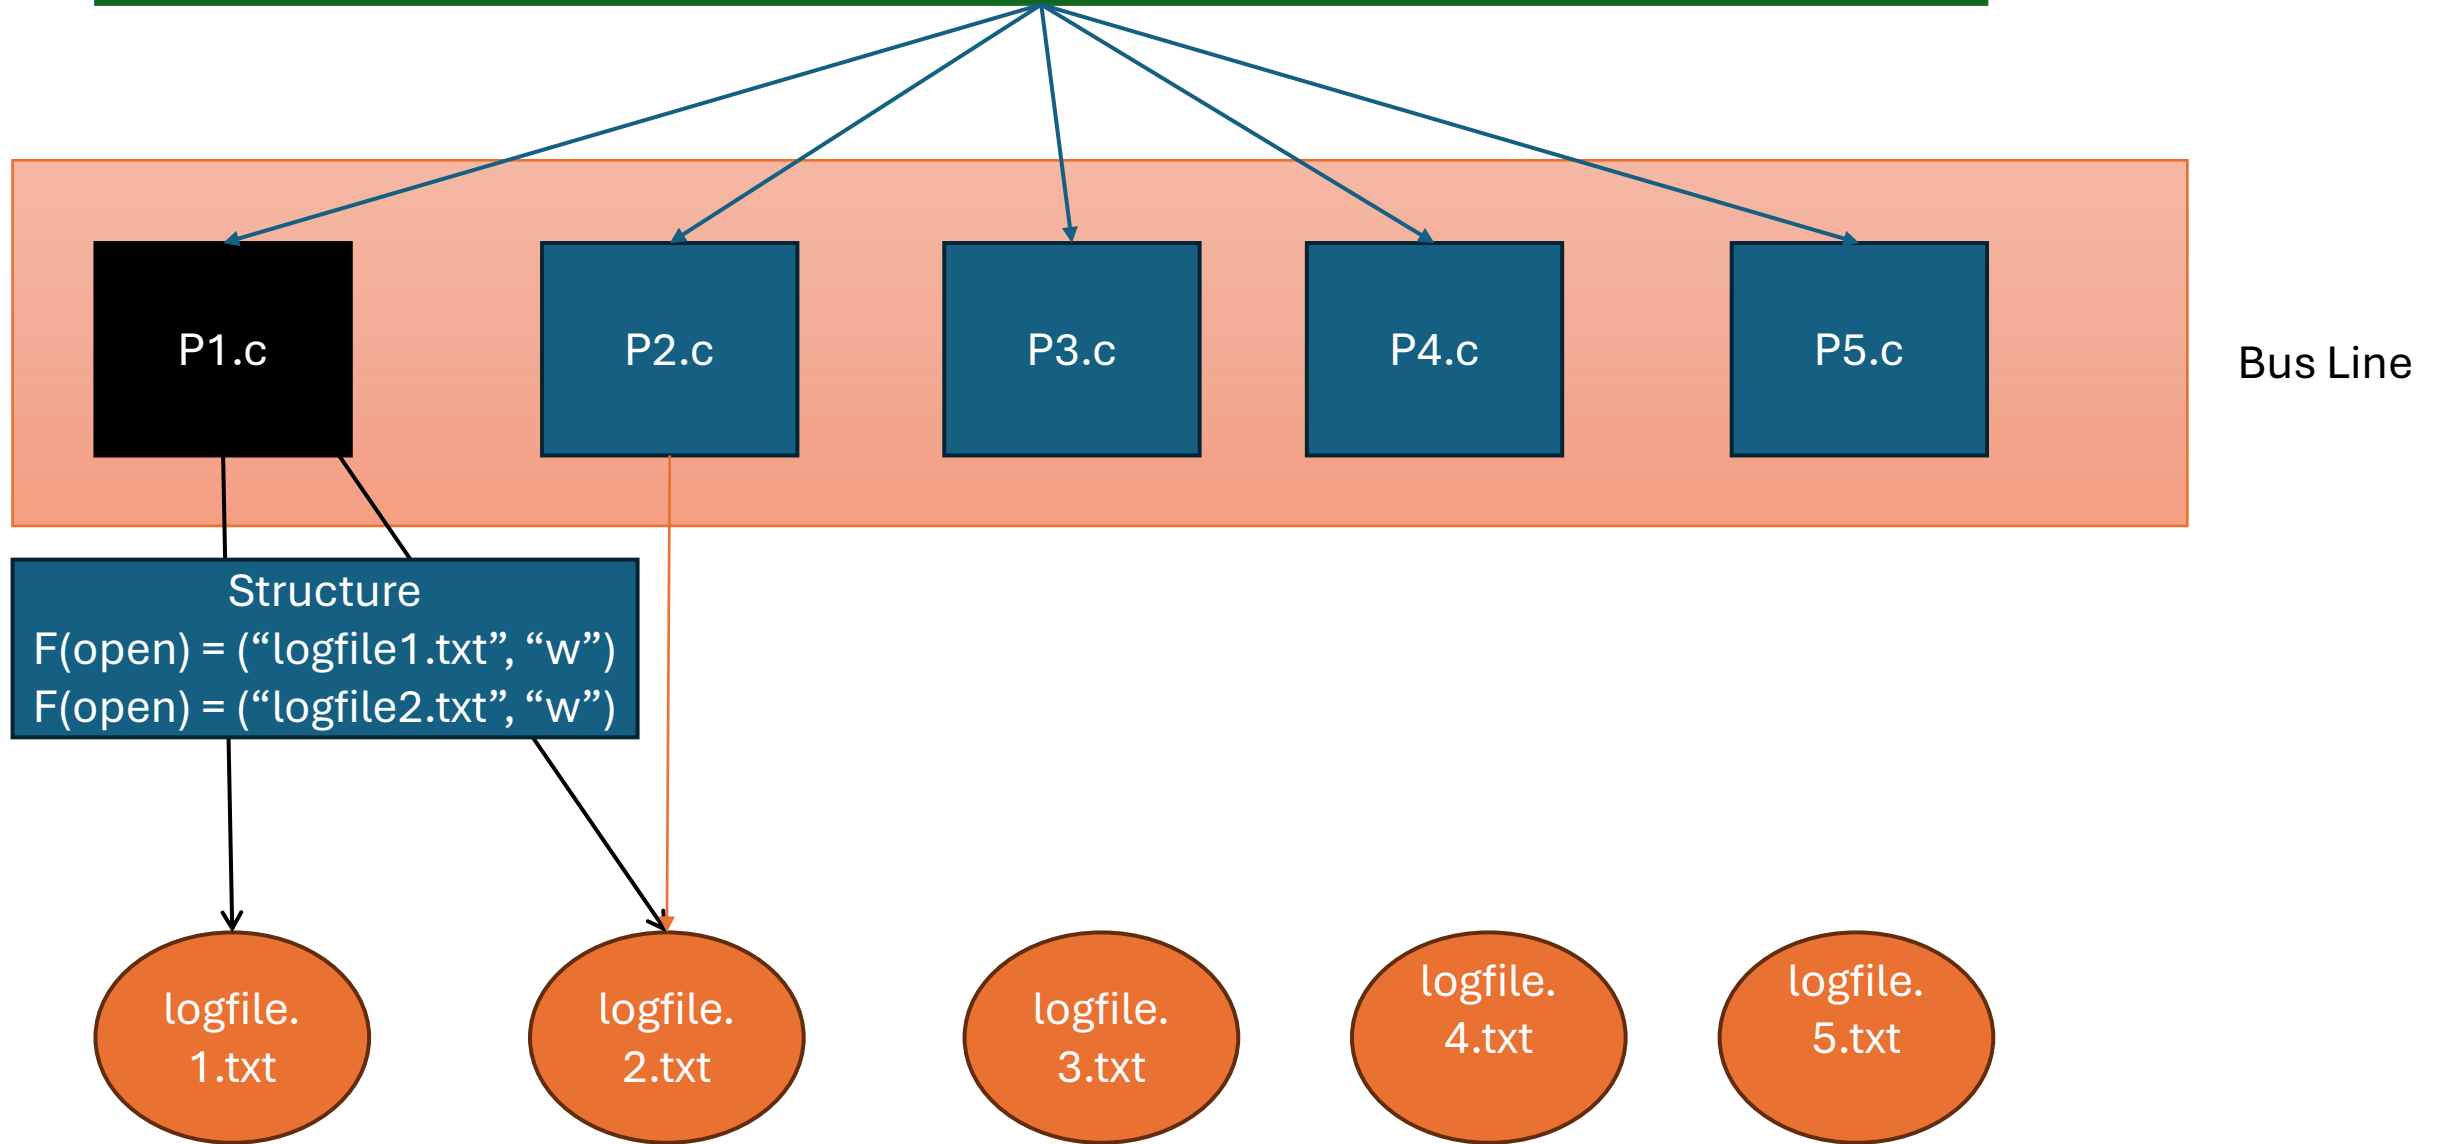
Shared Memory







Random Value Generator



What don't we know?

- > how many code files we need required?
- > I mean each process have each code file?
- >How different multiple process are communication each other?
- > How many processes need to take for instance?
- >If one more process is added, how it is scalable?

>How to identify which process came first?

